

Milestone 2 - StormCare RAG (Track C)

Track: Retrieval-Augmented Generation (Track C)

Tokenizer: distilgpt2 tokenizer (vocab size 50257)

Justification: matches generator, efficient for English disaster data

Data: FEMA, NOAA, CDC, Red Cross text (subset)

Pipeline: tokenization → dataloader → model → training loop

Model: DistilGPT2 causal LM

Training: AdamW, lr=5e-5, ~20 steps

Result: finite and stable loss, pipeline works end-to-end

Future: retrieval + RAG integration