

Milestone 2 Report - StormCare RAG (Track C)

Track: Retrieval-Augmented Generation (Track C)

Tokenizer: Using pretrained tokenizer (TinyLlama / DistilGPT2 compatible).

Reason: Matches generator model and avoids mismatch issues.

Data Pipeline: Text corpus from disaster preparedness sources (FEMA/NOAA style).

Context length: 128, Batch size: 2

Model: Small causal language model (smoke test only).

Training: 20 steps, AdamW optimizer.

Result: Loss is finite and stable. Pipeline works end-to-end.